

# Monthly PaNGalaxy Community Meeting

4 November 2022

ORNL is managed by UT-Battelle LLC for the US Department of Energy

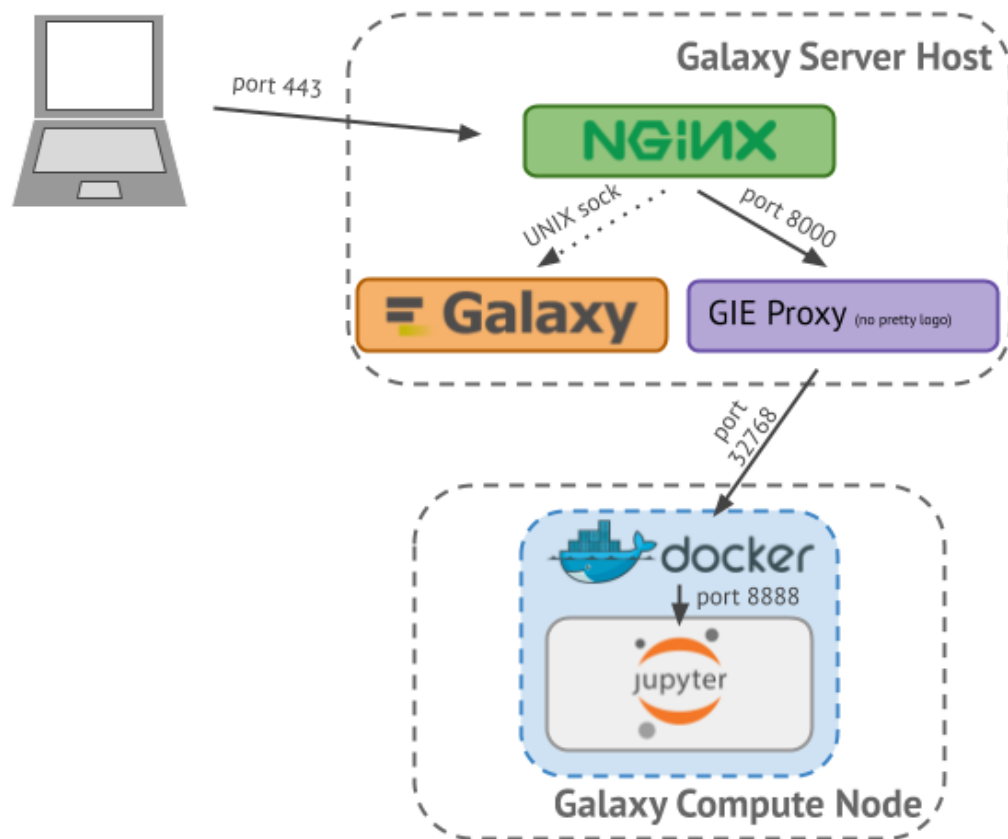


Research sponsored by the Laboratory Directed Research and Development Program of Oak Ridge National Laboratory, managed by UT- Battelle, LLC, for the U. S. Department of Energy

# Agenda

- Welcome for our colleagues from Muon Spectroscopy!
- Prototyping tools using Jupyter notebooks (ORNL)
- Deploying Galaxy on ORNL infrastructure (ORNL)
- Community topics:
  - Common tools/visualizations/datatypes
  - Event monitoring service
  - Muon Spectroscopy community
- Other topics
- Next meeting

# Jupyter notebooks – our experience




- We have only used Jupyter notebooks so far
  - Rstudio, other environments?
- Created own Docker image (Jupyter Lab + Mantid)
- Created own tools
  - generic one and a specific neutron tool
  - pushed to <https://github.com/PaNGalaxy/neutron-tools>
- Gained some experience in how to use them
  - no user feedback yet

<https://training.galaxyproject.org/training-material/topics/admin/tutorials/interactive-tools/tutorial.html>

# Jupyter notebooks – integration into a workflow

## Different modes

- Use as a standalone interactive tool
  - input/output via Bioblend
  - input/output via the tool
- Use as a standalone non-interactive tool
- Use as a part of a workflow
  - non-interactive
  - interactive?
- Reuse a notebook
- Create a regular tool out of a notebook

 **Interactive Jupyter Notebook ORNL** (Galaxy Version 0.1)

**Jupyter notebook**

Start with a fresh notebook

Select 'Start with a fresh notebook' to use a default one.

**Execute notebook noninteractively**


☒ No

This option is useful in workflows when you just want to execute a notebook and not

**Use dataset collection for input**


☒ No

**Datasets**

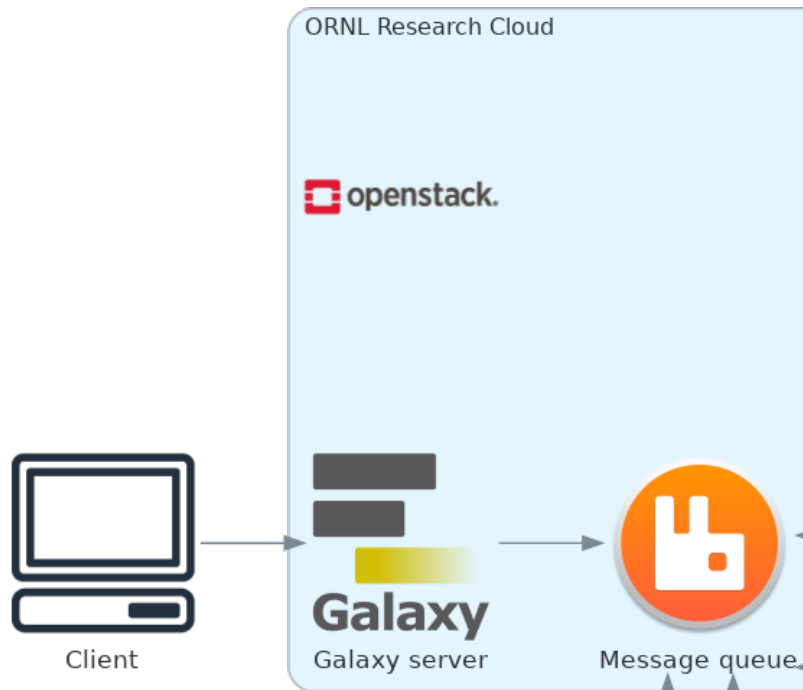
 Insert Datasets

**Type of output**

Single Dataset

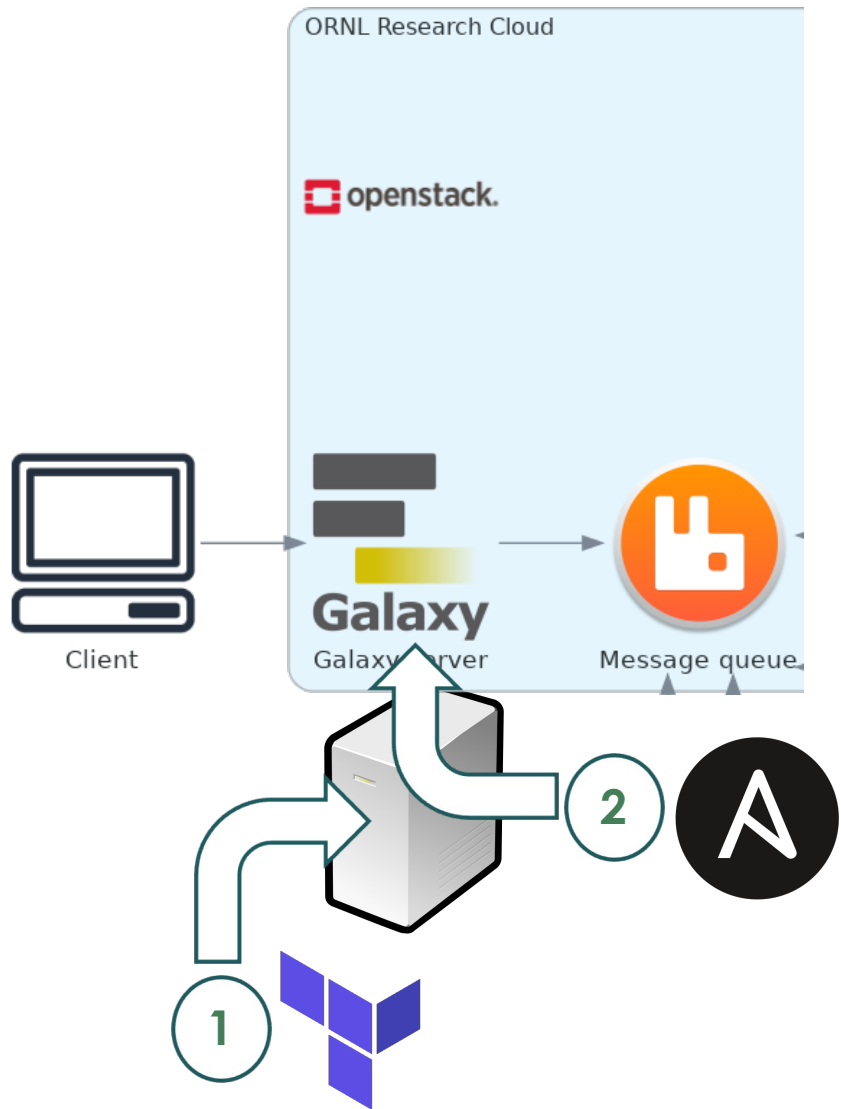
 Execute

# ORNL Neutrons Galaxy – Current Architecture



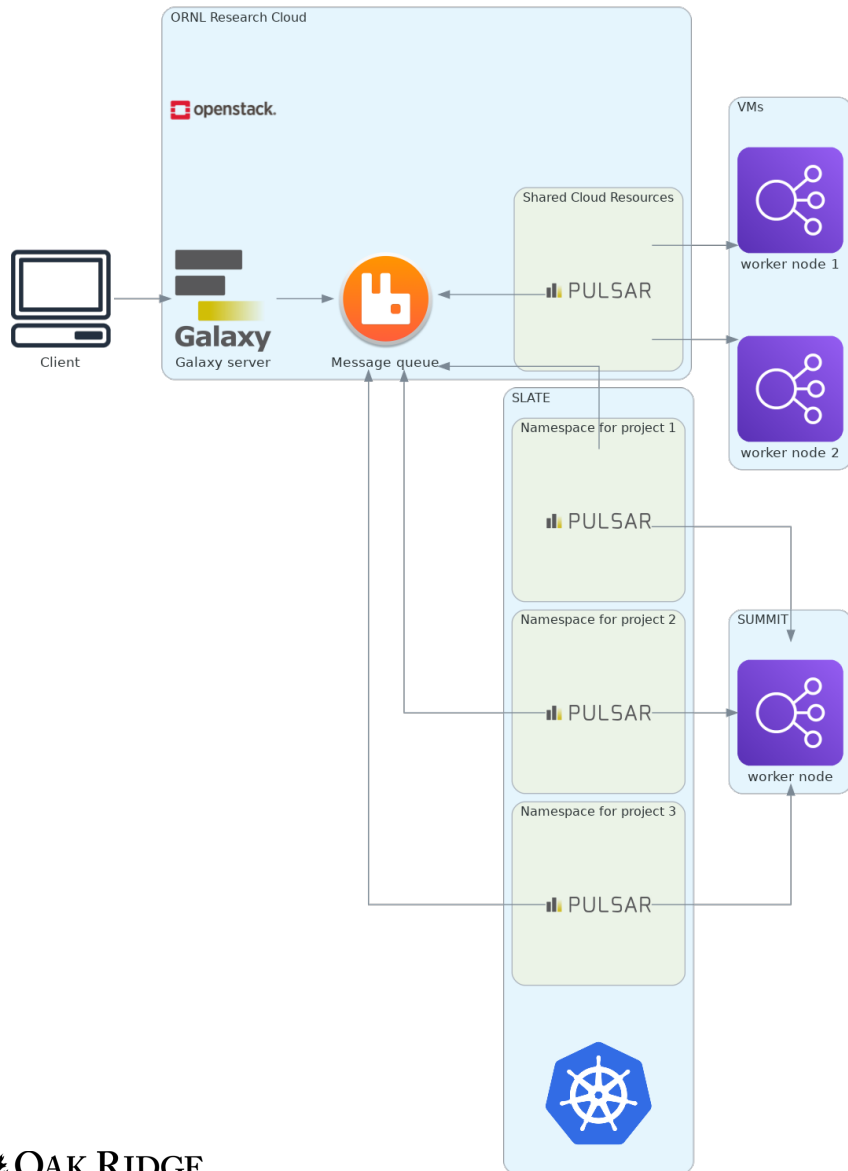
- ORNL Research Cloud
  - On-premise OpenStack cloud instance
- Galaxy Server
  - Virtual machine (VM) in ORNL Research Cloud
- Message Queue
  - RabbitMQ
  - VM in ORNL Research Cloud
  - Used to queue up work for compute resources

# ORNL Neutrons Galaxy – Current Deployment



1. Provision Infrastructure
  - Use Terraform, an Infrastructure-as-Code tool
  - Creates the cloud resource (i.e. VM + network rules / security group rules)
2. Configure / Install Galaxy
  - Configured via modified Galaxy Ansible playbooks
    - Examples:
      - <https://github.com/galaxyproject/ansible-galaxy>
      - <https://github.com/galaxyproject/ansible-nginx>
    - Modifications include...
      - Adding wildcard TLS/SSL certificates to server for the Jupyter Notebooks capabilities
      - Fixes for file permissions and ownership
      - Still others needed for more reliable installs!
  - Galaxy containers for Development
    - Maintain Galaxy containers we use for some development
    - Not currently used in our deployments

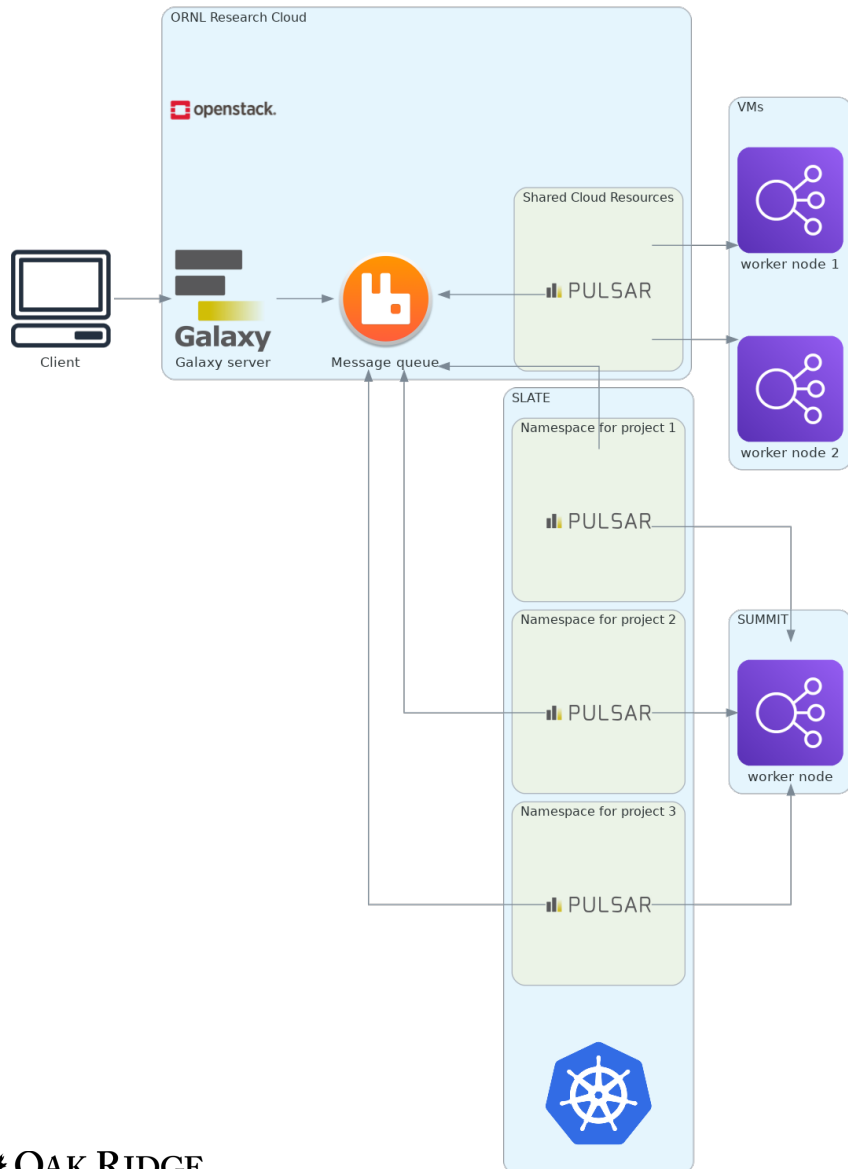
# ORNL Neutrons Galaxy – Planned Architecture Setup



- Pulsar
  - Manage work issued by Galaxy to compute resources
  - Matches requirements of hardware to parts of the workflows
- Compute in the Cloud
  - Easy setup, just VMs
  - CPU and CPU + GPU available in our cloud
  - Good for small to moderate jobs
- Compute on HPC
  - Slate Cluster
    - Kubernetes (well, OpenShift...) to host containers able to submit jobs to HPC resources (i.e. Summit for now, Frontier eventually)
    - Ingress / egress of commands and data to shared HPC User Facility
  - Summit
    - Where actual compute will run
    - Behind a job scheduler
  - Good for large jobs; User must setup the Pulsar containers in their Namespace of the Project



# ORNL Neutrons Galaxy – Future Deployment Discussions



- Containerization?
  - We use containers in development. Can we use those for deployment?
  - Have container orchestration setup in ORNL Research Cloud (i.e. Kubernetes clusters)
  - Unknowns:
    - Accessing the Neutrons Sciences NFS, which we don't own via a Pod in a Kubernetes cluster (persistent volume with the NFS?) Permissions?
    - Tools run on central server and spawn containers. Docker-in-docker scenario?
    - Jupyter Notebooks
      - Still able to spawn inside container?
      - Handling the wildcard certificate: How to set that up with ingress controllers and dynamic domain names for notebooks?
- Scaling?
  - Galaxy is a monolith in many ways
    - Frontend + backend
    - Central server is not distributed
- Monitoring?
  - Maybe Prometheus + Node Exporter (+ Grafana)